

Exercise10_ReckermannKuschWeinreich

January 16, 2025

1 Sheet 10

1.1 1 Optimal Transport

1.1

$$a) \min_p \sum_{i=1}^m \sum_{j=1}^m p_{ij} c_{ij} \quad \begin{array}{l} c_{ij} = \text{distances} \\ p_{ij} = \text{transport matrix} \end{array}$$

$$\text{capacity: } a_i = \sum_{j=1}^m p_{ij}$$

$$\text{demand: } b_j = \sum_{i=1}^m p_{ij} \quad p_{ij} \geq 0$$

c_{ij} has the smallest values if we transport on a straight line. With smaller values for c_{ij} the transport costs decrease.

$$b) \text{ we define: } \langle A, B \rangle = \sum_{i,j} A_{ij} B_{ij}$$

$$P \cdot \gamma_m = a \quad (\Leftrightarrow) \quad \begin{cases} P \gamma_m \geq a \\ P \gamma_m \leq a \end{cases} \quad (\Leftrightarrow) \quad \begin{bmatrix} P \\ -P \end{bmatrix} \gamma_m \geq \begin{bmatrix} a \\ -a \end{bmatrix}$$

$$\min_p \langle p, \gamma \rangle$$

$$\Rightarrow \begin{bmatrix} P \\ -P \end{bmatrix} \gamma_m \geq \begin{bmatrix} a \\ -a \end{bmatrix}, \quad \begin{bmatrix} P^T \\ -P^T \end{bmatrix} \gamma_m \geq \begin{bmatrix} b \\ -b \end{bmatrix}$$

```
[ ]: import numpy as np
import matplotlib.pyplot as plt

d = 5
num_sources = 10
num_sinks = 20
```

```

np.random.seed(42)
mass_sources = np.random.random(num_sources)
mass_sinks = np.random.random(num_sinks)
mass_sources /= np.sum(mass_sources)
mass_sinks /= np.sum(mass_sinks)

coords_sources = np.random.rand(num_sources, d)
coords_sinks = np.random.rand(num_sinks, d)

```

```

[26]: # TODO: solve the OT problem as linear program

distances = np.zeros((num_sources, num_sinks))
for i in range(num_sources):
    for j in range(num_sinks):
        distances[i, j] = np.linalg.norm(coords_sources[i] - coords_sinks[j])

distances.shape

from scipy.optimize import linprog

c = distances.flatten()

A_eq = np.zeros((num_sources + num_sinks, num_sources * num_sinks))
b_eq = mass_sources.tolist() + mass_sinks.tolist()

for i in range(num_sources):
    for j in range(num_sinks):
        A_eq[i, i*num_sinks + j] = 1

for i in range(num_sinks):
    for j in range(num_sources):
        A_eq[num_sources + i, j*num_sinks + i] = 1

res = linprog(c, A_eq=A_eq, b_eq=b_eq, bounds=(0, None))

cost = c.T@res.x
print(f"Transport cost: {cost:.3f}")

```

Transport cost: 0.671

Very good

1.1.1 d)

Let the intersecting routes transport T_{ij} from $x_i \rightarrow y_j$ and T_{kl} from $x_k \rightarrow y_l$. The total cost of these two routes is:

$$C_{\text{intersect}} = T_{ij} \cdot \|x_i - y_j\| + T_{kl} \cdot \|x_k - y_l\|.$$

Now, consider rerouting the transportation: - Transport T_{ij} from $x_i \rightarrow y_l$. - Transport T_{kl} from

$x_k \rightarrow y_j$.

The new total cost is:

$$C_{\text{reroute}} = T_{ij} \cdot \|x_i - y_l\| + T_{kl} \cdot \|x_k - y_j\|.$$

Using the triangle inequality:

What about settings where sinks and sources don't all match up?

$$\|x_i - y_j\| + \|x_k - y_l\| \geq \|x_i - y_l\| + \|x_k - y_j\|.$$

Equality holds only if x_i, y_j, x_k, y_l lie on a straight line in the order of transportation. Otherwise:

$$C_{\text{intersect}} > C_{\text{reroute}}.$$

Since rerouting reduces the cost, any transportation plan with intersecting routes is suboptimal. Therefore, optimal transportation plans do not have intersecting routes.

1.2 2 Flow matching for generative modeling

```
[ ]: import torch

def generate_checkerboard_sample(num_samples=10, field_size=0.4, num_fields=2,
    center=True):
    x = torch.rand(num_samples, 2) * field_size
    offset = torch.randint(0, num_fields, (num_samples, 2)) * field_size * 2
    diagonal_shift = torch.randint(0, num_fields, (num_samples, 1)) * field_size
    x += offset + diagonal_shift

    if center:
        x -= torch.mean(x, dim=0)

    return x

base_distribution_std = 0.15
num_samples = 2000
x = torch.randn(num_samples, 2) * base_distribution_std
y = generate_checkerboard_sample(num_samples=num_samples)

# show points
plt.scatter(x[:, 0], x[:, 1], alpha=0.5, label='base distribution')
plt.scatter(y[:, 0], y[:, 1], alpha=0.5, label='checkerboard distribution')
plt.show()
```

```
[ ]: # define a model
from torchvision.ops import MLP
from tqdm import tqdm

device = "cuda" if torch.cuda.is_available() else "cpu"
```



```

model = MLP(in_channels=2 + 1, hidden_channels=[512, 512, 512, 512, 2],
    activation_layer=torch.nn.SiLU)
model.to(device)

# define a loss function
criterion = torch.nn.MSELoss(reduction="none")

# define an optimizer
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)

# train the model:
num_epochs = 20000 # use fewer epochs if it takes too long
batch_size = 4096
losses = []
for epoch in tqdm(range(num_epochs)):
    x = torch.randn(batch_size, 2) * base_distribution_std
    y = generate_checkerboard_sample(num_samples=batch_size)

    # TODO: implement the training loop

```

```

[23]: # TODO: run inference with the trained model.
# Visualize the trajectory of the samples and the final samples at t=1.
# Hint: Use a simple Euler integration scheme to integrate the velocity field
    with 100 steps.

```

1.3 3 Adversarial attacks and AI safety

```

[24]: '''
Logistic regression in pytorch (needed for backpropagation)
taken from https://github.com/saprmrks/geometry-of-truth/blob/main/probes.py
'''

class LRProbe(torch.nn.Module):
    def __init__(self, d_in):
        super().__init__()
        self.net = torch.nn.Sequential(
            torch.nn.Linear(d_in, 1, bias=False),
            torch.nn.Sigmoid()
        )

    def forward(self, x):
        return self.net(x).squeeze(-1)

    def pred(self, x):
        return self(x).round()

```

```

    def from_data(acts, labels, lr=0.001, weight_decay=0.1, epochs=1000,
device='cpu'):
        acts, labels = acts.to(device), labels.to(device)
        probe = LRProbe(acts.shape[-1]).to(device)

        opt = torch.optim.AdamW(probe.parameters(), lr=lr,
weight_decay=weight_decay)
        for _ in range(epochs):
            opt.zero_grad()
            loss = torch.nn.BCELoss()(probe(acts), labels)
            loss.backward()
            opt.step()

        return probe

    def __str__():
        return "LRProbe"

    @property
    def direction(self):
        return self.net[0].weight.data[0]

```

```

[25]: # We import the DataManager class as a helper function to load the activation
vectors for us.
from lie_detection_utils import DataManager
from sklearn.metrics import accuracy_score

path_to_datasets = "data/lie_detection/datasets"
path_to_acts = "data/lie_detection/acts"

```

```

[ ]: # train a model on the cities dataset
dataset_name = "cities"

dm = DataManager()
dm.add_dataset(dataset_name, "Llama3", "8B", "chat", layer=12, split=0.8,
center=False,
device='cpu', path_to_datasets=path_to_datasets,
path_to_acts=path_to_acts)
train_acts, train_labels = dm.get('train')
test_acts, test_labels = dm.get('val')

print("train_acts.shape", train_acts.shape)
print("test_acts.shape", test_acts.shape)

# TODO: train a logistic regression probe on the train_acts and train_labels

```

```
[ ]: # TODO: optimize a perturbation on a single sample which is a lie
[ ]: # TODO: check whether this perturbation works on other samples too
[33]: # TODO: add the constraint that the perturbation should be small
```